

```

# Simplified TensorFlow with Gym for Google Colab - Faster Version
# May 2025

import tensorflow as tf
import gymnasium as gym
import numpy as np
import matplotlib.pyplot as plt

# Install required packages (uncomment if needed)
# !pip install gymnasium

# Create environment
env = gym.make('CartPole-v1')
state_size = env.observation_space.shape[0] # 4 values: position,
velocity, angle, angular velocity
action_size = env.action_space.n           # 2 actions: move left or
right

# Build a smaller model
model = tf.keras.Sequential([
    tf.keras.layers.Dense(12, input_dim=state_size,
activation='relu'),
    tf.keras.layers.Dense(action_size, activation='linear')
])
model.compile(loss='mse',
optimizer=tf.keras.optimizers.Adam(learning_rate=0.001))
model.summary()

# Training parameters - REDUCED FOR SPEED
epsilon = 1.0           # Exploration rate
epsilon_decay = 0.9     # Faster decay
epsilon_min = 0.1
gamma = 0.95            # Discount factor
memory = []
batch_size = 16         # Smaller batch
episodes = 10           # Fewer episodes

# Track scores for plotting
scores = []

# Training loop
for episode in range(episodes):
    # Reset environment
    state, _ = env.reset()
    state = np.reshape(state, [1, state_size])
    done = False
    score = 0

    # Limit steps per episode for speed
    max_steps = 200

```

```

for step in range(max_steps):
    # Choose action: explore or exploit
    if np.random.rand() <= epsilon:
        action = env.action_space.sample() # Explore
    else:
        action_values = model.predict(state, verbose=0)
        action = np.argmax(action_values[0]) # Exploit

    # Take action
    next_state, reward, terminated, truncated, _ =
env.step(action)
    done = terminated or truncated
    next_state = np.reshape(next_state, [1, state_size])
    score += reward

    # Store experience
    memory.append((state, action, reward, next_state, done))
    state = next_state

    # Train model - only train every 4 steps for speed
    if len(memory) > batch_size and step % 4 == 0:
        # Sample random batch from memory
        batch = np.random.choice(len(memory), batch_size,
replace=False)
        for i in batch:
            s, a, r, next_s, d = memory[i]
            # Q-learning formula
            target = r
            if not d:
                target = r + gamma * np.amax(model.predict(next_s,
verbose=0)[0])

            target_f = model.predict(s, verbose=0)
            target_f[0][a] = target

            # Train network with a single sample
            model.fit(s, target_f, epochs=1, verbose=0)

        if done:
            break

    # Reduce exploration faster
    epsilon = max(epsilon_min, epsilon * epsilon_decay)

    scores.append(score)
    print(f"Episode: {episode+1}/{episodes}, Score: {score}, Epsilon:
{epsilon:.2f}")

# Plot results
plt.figure(figsize=(8, 4))

```

```

plt.plot(scores)
plt.title('Training Progress')
plt.xlabel('Episodes')
plt.ylabel('Score')
plt.grid(True)
plt.show()

# Quick test of trained model
print("\nTesting trained model...")
state, _ = env.reset()
state = np.reshape(state, [1, state_size])
done = False
score = 0
for _ in range(200): # Limit test steps
    action_values = model.predict(state, verbose=0)
    action = np.argmax(action_values[0])
    next_state, reward, terminated, truncated, _ = env.step(action)
    done = terminated or truncated
    state = np.reshape(next_state, [1, state_size])
    score += reward
    if done:
        break

print(f"Test score: {score}")
env.close()

```

Model: "sequential\_3"

| Layer (type)<br>Param # | Output Shape |  |
|-------------------------|--------------|--|
| dense_12 (Dense)<br>60  | (None, 12)   |  |
| dense_13 (Dense)<br>26  | (None, 2)    |  |

Total params: 86 (344.00 B)

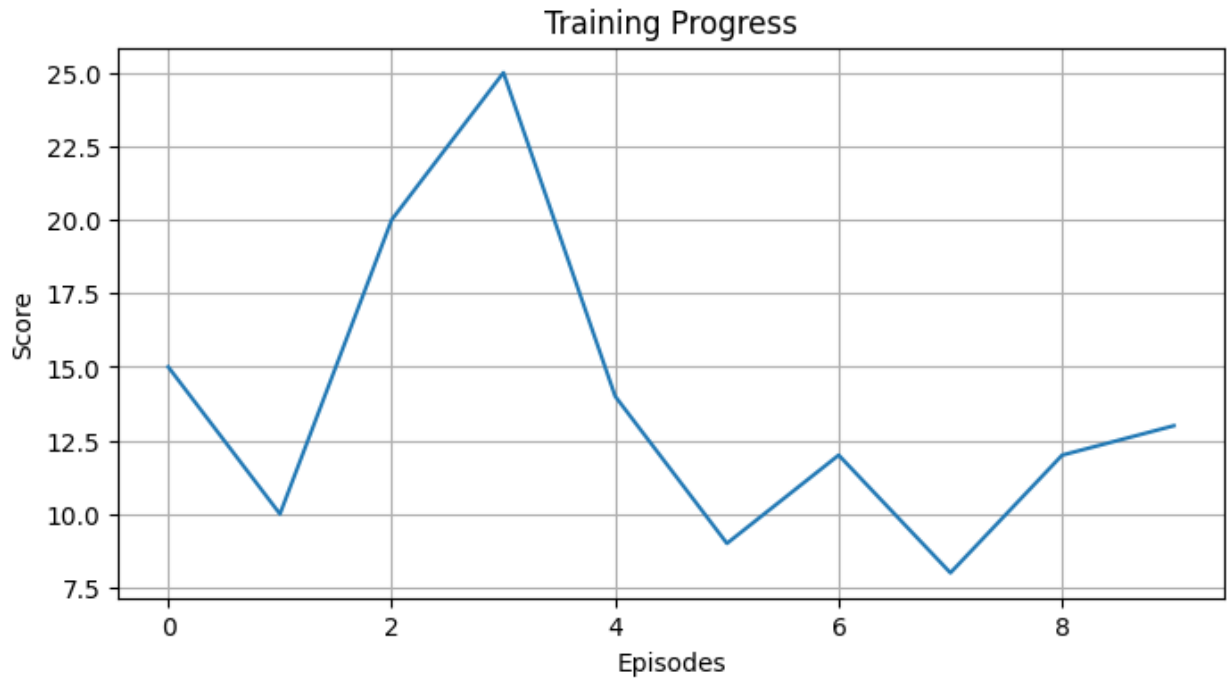
Trainable params: 86 (344.00 B)

Non-trainable params: 0 (0.00 B)

Episode: 1/10, Score: 15.0, Epsilon: 0.90

Episode: 2/10, Score: 10.0, Epsilon: 0.81

Episode: 3/10, Score: 20.0, Epsilon: 0.73  
Episode: 4/10, Score: 25.0, Epsilon: 0.66  
Episode: 5/10, Score: 14.0, Epsilon: 0.59  
Episode: 6/10, Score: 9.0, Epsilon: 0.53  
Episode: 7/10, Score: 12.0, Epsilon: 0.48  
Episode: 8/10, Score: 8.0, Epsilon: 0.43  
Episode: 9/10, Score: 12.0, Epsilon: 0.39  
Episode: 10/10, Score: 13.0, Epsilon: 0.35



Testing trained model...  
Test score: 10.0