

```

import gymnasium as gym
import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import matplotlib.pyplot as plt
# Hyperparameters
EPISODES = 50          # Increased slightly for training signal
MAX_STEPS = 200
GAMMA = 0.99
LR = 0.001
ENTROPY_BETA = 0.01    # Encourage exploration

# Environment
env = gym.make("Acrobot-v1")
state_dim = env.observation_space.shape[0]
action_dim = env.action_space.n

# Actor-Critic model
class ActorCritic(nn.Module):
    def __init__(self):
        super().__init__()
        self.shared = nn.Sequential(
            nn.Linear(state_dim, 128),
            nn.ReLU()
        )
        self.policy = nn.Sequential(
            nn.Linear(128, action_dim),
            nn.Softmax(dim=-1)
        )
        self.value = nn.Linear(128, 1)

    def forward(self, x):
        x = self.shared(x)
        return self.policy(x), self.value(x)

model = ActorCritic()
optimizer = optim.Adam(model.parameters(), lr=LR)

# Training loop
all_rewards = []

for episode in range(EPISODES):
    state, _ = env.reset()
    log_probs = []
    values = []
    rewards = []
    entropies = []
    total_reward = 0

```

```

for _ in range(MAX_STEPS):
    state_tensor = torch.FloatTensor(state).unsqueeze(0)
    probs, value = model(state_tensor)
    dist = torch.distributions.Categorical(probs)
    action = dist.sample()
    log_prob = dist.log_prob(action)
    entropy = dist.entropy()

    next_state, reward, terminated, truncated, _ =
env.step(action.item())
    done = terminated or truncated

    # Reward shaping: add small bonus if the tip is closer to top
    reward += 1.0 * (abs(state[0]) + abs(state[1])) # cosine of
angles closer to 1 = higher

    log_probs.append(log_prob)
    values.append(value)
    rewards.append(reward)
    entropies.append(entropy)
    total_reward += reward

    state = next_state
    if done:
        break

# Compute returns and advantage
    returns = []
    R = 0
    for r in reversed(rewards):
        R = r + GAMMA * R
        returns.insert(0, R)
    returns = torch.FloatTensor(returns)
    values = torch.cat(values)
    log_probs = torch.stack(log_probs)
    entropies = torch.stack(entropies)
    advantage = returns - values.squeeze()

# Loss
    actor_loss = -(log_probs * advantage.detach()).mean()
    critic_loss = advantage.pow(2).mean()
    entropy_loss = -ENTROPY_BETA * entropies.mean()
    loss = actor_loss + 0.5 * critic_loss + entropy_loss

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    all_rewards.append(total_reward)

```

```
print(f"Ep {episode+1}/{EPISODES}, Total Reward:  
{total_reward:.1f}")
```

```
Ep 1/50, Total Reward: 28.7  
Ep 2/50, Total Reward: 54.9  
Ep 3/50, Total Reward: 45.7  
Ep 4/50, Total Reward: 49.9  
Ep 5/50, Total Reward: 50.1  
Ep 6/50, Total Reward: 56.9  
Ep 7/50, Total Reward: 26.8  
Ep 8/50, Total Reward: 56.1  
Ep 9/50, Total Reward: 32.0  
Ep 10/50, Total Reward: 38.7  
Ep 11/50, Total Reward: 51.8  
Ep 12/50, Total Reward: 35.4  
Ep 13/50, Total Reward: 27.5  
Ep 14/50, Total Reward: 44.4  
Ep 15/50, Total Reward: 48.6  
Ep 16/50, Total Reward: 26.7  
Ep 17/50, Total Reward: 39.1  
Ep 18/50, Total Reward: 35.2  
Ep 19/50, Total Reward: 50.4  
Ep 20/50, Total Reward: 52.2  
Ep 21/50, Total Reward: 48.6  
Ep 22/50, Total Reward: 48.6  
Ep 23/50, Total Reward: 41.5  
Ep 24/50, Total Reward: 27.9  
Ep 25/50, Total Reward: 54.7  
Ep 26/50, Total Reward: 35.9  
Ep 27/50, Total Reward: 44.4  
Ep 28/50, Total Reward: 44.8  
Ep 29/50, Total Reward: 42.7  
Ep 30/50, Total Reward: 37.0  
Ep 31/50, Total Reward: 39.3  
Ep 32/50, Total Reward: 27.4  
Ep 33/50, Total Reward: 41.7  
Ep 34/50, Total Reward: 22.8  
Ep 35/50, Total Reward: 48.8  
Ep 36/50, Total Reward: 40.4  
Ep 37/50, Total Reward: 36.6  
Ep 38/50, Total Reward: 43.6  
Ep 39/50, Total Reward: 53.0  
Ep 40/50, Total Reward: 50.7  
Ep 41/50, Total Reward: 25.2  
Ep 42/50, Total Reward: 55.3  
Ep 43/50, Total Reward: 29.7  
Ep 44/50, Total Reward: 45.2  
Ep 45/50, Total Reward: 31.3  
Ep 46/50, Total Reward: 51.3  
Ep 47/50, Total Reward: 38.8
```

```
Ep 48/50, Total Reward: 50.4  
Ep 49/50, Total Reward: 45.3  
Ep 50/50, Total Reward: 48.3
```

```
# Plot results  
plt.plot(all_rewards)  
plt.xlabel("Episode")  
plt.ylabel("Total Reward (Shaped)")  
plt.title("Actor-Critic on Acrobot-v1")  
plt.grid(True)  
plt.show()
```

